

Linear and Logistic Regression

Simple, yet powerful predictors

Yordan Darakchiev

Technical Trainer

iordan93@gmail.com



Table of Contents

- sli.do: [#regr-models](#)
- Machine learning basics
 - Objective function, cost function, optimization
- Linear regression
 - Problem description, motivation
 - Algorithm
 - Usage
- RANSAC
- Extensions: polynomial regression
- Logistic regression
 - Problem description
 - Algorithm
 - Usage

Linear Regression


**Predict continuous values...
and torture first-semester students**

Linear Regression Intuition

- Regression – predicting a continuous variable
- Problem statement
 - Given pairs of $(x; y)$ points, create a model
 - Input x , output y ; **goal: predict y given x**
 - Under the assumption that y depends linearly on x (and nothing else)
- Linear regression model
 - $y = ax + b$
 - **Modelling function**
 - a, b – unknown parameters
 - Example: $y = 2x + 3$
 - Real case: we have many sources of error
 - So, the relationship we observe, cannot be perfect
 - There is some noise added to our data
 - $y = ax + b + \varepsilon$, ε – noise
 - We **don't want** to model the noise, only the “useful function”


Loss Function

- Consider the distances from each point to the line
 - Vertical distances: $d = y - \tilde{y}$
 - Better measure: squared distance (always positive)
 - \tilde{y} – predicted value: $\tilde{y} = ax + b$
- Total distance: $\sum d_i$
 - Lower total distance = better model
- **Total loss function:** $J = \frac{1}{n} \sum_{i=1}^n (y_i - \tilde{y}_i)^2$
 - Smaller value \Rightarrow better model
 - Objective: find $\min_{a,b} J$



Gradient Descent

- Input: a, b ; output J
- Paraboloid (3D parabola)
 - It has exactly one min value
 - And we can see it
- Intuition
 - If the plot was a real object (say, a sheet of some sort), we could slide a ball bearing on it
 - After a while, the ball bearing will settle at the “bottom” due to gravity
 - We can "simulate" this: **gradient descent**
- Reminder: gradient
 - "Multi-dimensional derivative"



$$\nabla J = \begin{pmatrix} \frac{\partial J}{\partial a} \\ \frac{\partial J}{\partial b} \end{pmatrix}$$

Gradient Descent (2)


- Iterative algorithm – perform as long as needed
 - Start from some point in the $(a; b)$ space: $(a_0; b_0)$
 - Decide how big steps to take: number α
 - Called “learning rate” in ML terminology
 - Use the current a, b and x, y to compute ∇J
 - $-\nabla J_a$ tells us how much to move in the a direction in order to get to the minimum
 - Similar for $-\nabla J_b$
 - Take a step with size α in each direction
 - $a_1 = a_0 - \nabla J_a; b_1 = b_0 - \nabla J_b$
 - $(a_1; b_1)$ are the new coordinates
 - Repeat the two preceding steps as needed
 - Usually, we do this for a fixed number of iterations

Results and Interpretation

- Going through the entire process, we now have a line $\tilde{y} = ax + b$ which describes our data in the best way
 - We could plot the evolution of J to see that it always decreases
 - If it doesn't, this indicates a problem with our algorithm
- This was a lot of work
 - Thankfully, there are libraries that hide away all that complexity for us
 - `scikit-learn` is the most popular of them
 - Arguably, the most popular of the `scikits` as well
 - Also, generalizes trivially to more dimensions

```
from sklearn.linear_model import LinearRegression

model = LinearRegression()
model.fit(data_x.reshape(-1, 1), data_y)
print(model.coef_, model.intercept_)
```



Lab: Linear Regression on Real Data

- The algorithm can be generalized to more than 2D
 - "Multiple linear regression": $y = \beta_0 + \beta_1 x_1 + \dots + \beta_m x_m$
- Let's use this model to try and predict housing prices (a classical dataset located [here](#))

```
housing.columns = ["crime_rate", "zoned_land", "industry", "bounds_river",  
"nox_conc", "rooms", "age", "distance", "highways", "tax", "pt_ratio",  
"b_estimator", "pop_status", "price"]
```

- First, we want to explore the datasets
 - A more thorough exploration is "left as an exercise to the reader"
 - But we want to see what model would be appropriate
 - In addition to usual data analysis techniques, let's plot all correlations between any pair of features

Creating a Model

- Modelling is very simple
 - Like in the 2D example


```
housing_model = LinearRegression()  
predictor_attributes = housing.drop("price", axis = 1)  
housing_model.fit(predictor_attributes, housing.price)  
print(housing_model.coef_)  
print(housing_model.intercept_)
```

- So what?
 - We might want to predict some prices
 - Let's just pass some random rows and see the result
 - **Note: Never test on the training dataset!**

```
test_houses = housing.sample(10)  
predicted = housing_model.predict(  
    test_houses.drop("price", axis = 1))  
print(predicted)  
print(test_houses.price)
```

Regression with Outliers

- As we saw, the data has outliers
 - A few points which are far from the others
- Our goal is to exclude outliers
 - There are several methods
 - One very common – RANSAC (**RAN**dom **SA**mples **C**onsensus)
- Algorithm
 1. Fit a model to a random subsample ("inliers")
 2. Test all data points and include those which are "near" the model
 - Small enough error, tolerance provided by developer
 3. Fit the model again
 4. Estimate the error of the model (difference between first and second)
 5. Iterate steps 1-4 until performance reaches a threshold or number of iterations



Lab: RANSAC on the Housing Dataset

- Usage: similar to the linear regression model

```
from sklearn.linear_model import RANSACRegressor
ransac = RANSACRegressor()
ransac.fit(housing.drop("price", axis = 1), housing.price)
print(ransac.estimator_.coef_, ransac.estimator_.intercept_)
```

- We can also provide parameters, e.g. min number of random samples, max iterations, threshold (to include data points)
 - We can also provide the type of model we want to perform RANSAC on
 - Linear regression by default but we may use other regression models

```
ransac = RANSACRegressor(LinearRegression(), min_samples = 50,
    max_trials = 100, residual_threshold = 5.0)
```

- View inliers and outliers

```
inliers = housing[ransac.inlier_mask_]
outliers = housing[~ransac.inlier_mask_]
plt.scatter(inliers.rooms, inliers.price)
plt.scatter(outliers.rooms, outliers.price)
```

Polynomial Regression


- Extension of the linear regression algorithm
 - We can use the linear regression algorithm to perform polynomial regression (e.g. fitting a quadratic curve)
 - Just precompute the columns
 - Example: if we have columns x , y and z , compute $x * z$, $y * z$, $x * x$ and perform linear regression on these 6 features
 - Example 2: polynomial terms: multiply x by itself: $x * x$, $x * x * x$, etc.
- This can be achieved easily with `scikit-learn`

```
from sklearn.preprocessing import PolynomialFeatures


x = np.arange(6).reshape(3, 2)
poly = PolynomialFeatures(2)
x_transformed = poly.fit_transform(x)
print(poly.get_feature_names())
print(poly.n_input_features_)
print(poly.n_output_features_)
# Now we can perform linear regression with x_transformed as the input
```

Common Mistakes

- There are two main types of errors we can make while trying regression models
 - Use a **wrong model**
 - Anscombe's quartet
 - **Extrapolate** without knowing (especially if we have interacting features)



MY HOBBY: EXTRAPOLATING



Logistic Regression


Use a regression model to classify

Classification

- Predict **one of several known classes**
 - Based on the input parameters
 - Example: classify whether a picture is of a cat or a dog
- Regression and classification make up most of the machine learning problems
- Choosing an algorithm
 - "No free lunch": **no single algorithm** works best
 - It's best to compare some algorithms to select the best for a particular model
 - Also, we might want to tune them first
- Reminder: ML process
 - Select features, choose a performance metric (cost function), choose a classifier, evaluate and fine-tune the performance


Logistic Regression

- Classification algorithm (despite its name)
- Two classes: negative (0) and positive (1)
 - Can be extended to more classes
- How does it work?
 - Linear regression can give us all kinds of values
 - We want to constrain them between 0 and 1
 - Approach
 - Perform linear regression: $\tilde{y} = \beta x$
 - Use the sigmoid function to constrain the output:
$$\sigma(\tilde{y}) = \frac{1}{1 + e^{-\tilde{y}}} = \frac{1}{1 + e^{-\beta x}}$$
 - Quantization: if $\sigma > 0.5$ return 1, and 0 otherwise
 - Remember that we only need to return 0 or 1
 - We can also use the raw values as probability measures



Lab: Logistic Regression on Real Data

- A classic dataset for classification is the Iris dataset
 - Located [here](#)
 - **3 classes** (setosa, virginica, versicolor)
 - **4 attributes**: petal width / height; sepal width / height (all in cm)
 - Some features are highly correlated to the class
 - Explore and inspect the data before modelling



Lab: Logistic Regression on Real Data (2)

- Perform logistic regression

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression(C = 1e6)
model.fit(iris_train_data, iris_train_labels)
```

- Test (output classes or probabilities)

```
print(model.predict(iris_test))
print(model.predict_proba(iris_test))
```

- In the model, there's a "mysterious" parameter C
 - Regularization: how powerful the data is (more – next time)
 - A large number means no regularization
 - We just take the data "as-is", with no other constraints

Many Classes

- Two main approaches
 - One-vs-all: several predictors
 - One predictor for each class vs. the others
 - Overall: calculate probabilities of each class
- `scikit-learn` takes care of multiple classes (multinomial logistic regression) by default
 - We don't even need to transform the labels
 - This applies to all algorithms in the library

Summary

- Machine learning basics
 - Objective function, cost function, optimization
- Linear regression
 - Problem description, motivation
 - Algorithm
 - Usage
- RANSAC
- Extensions: polynomial regression
- Logistic regression
 - Problem description
 - Algorithm
 - Usage

The image features a white background with two blue decorative bars. The top bar is a solid blue strip. The bottom bar is a gradient blue strip that transitions from a lighter blue on the left to a darker blue on the right. The word "Questions?" is centered in a blue, sans-serif font.

Questions?